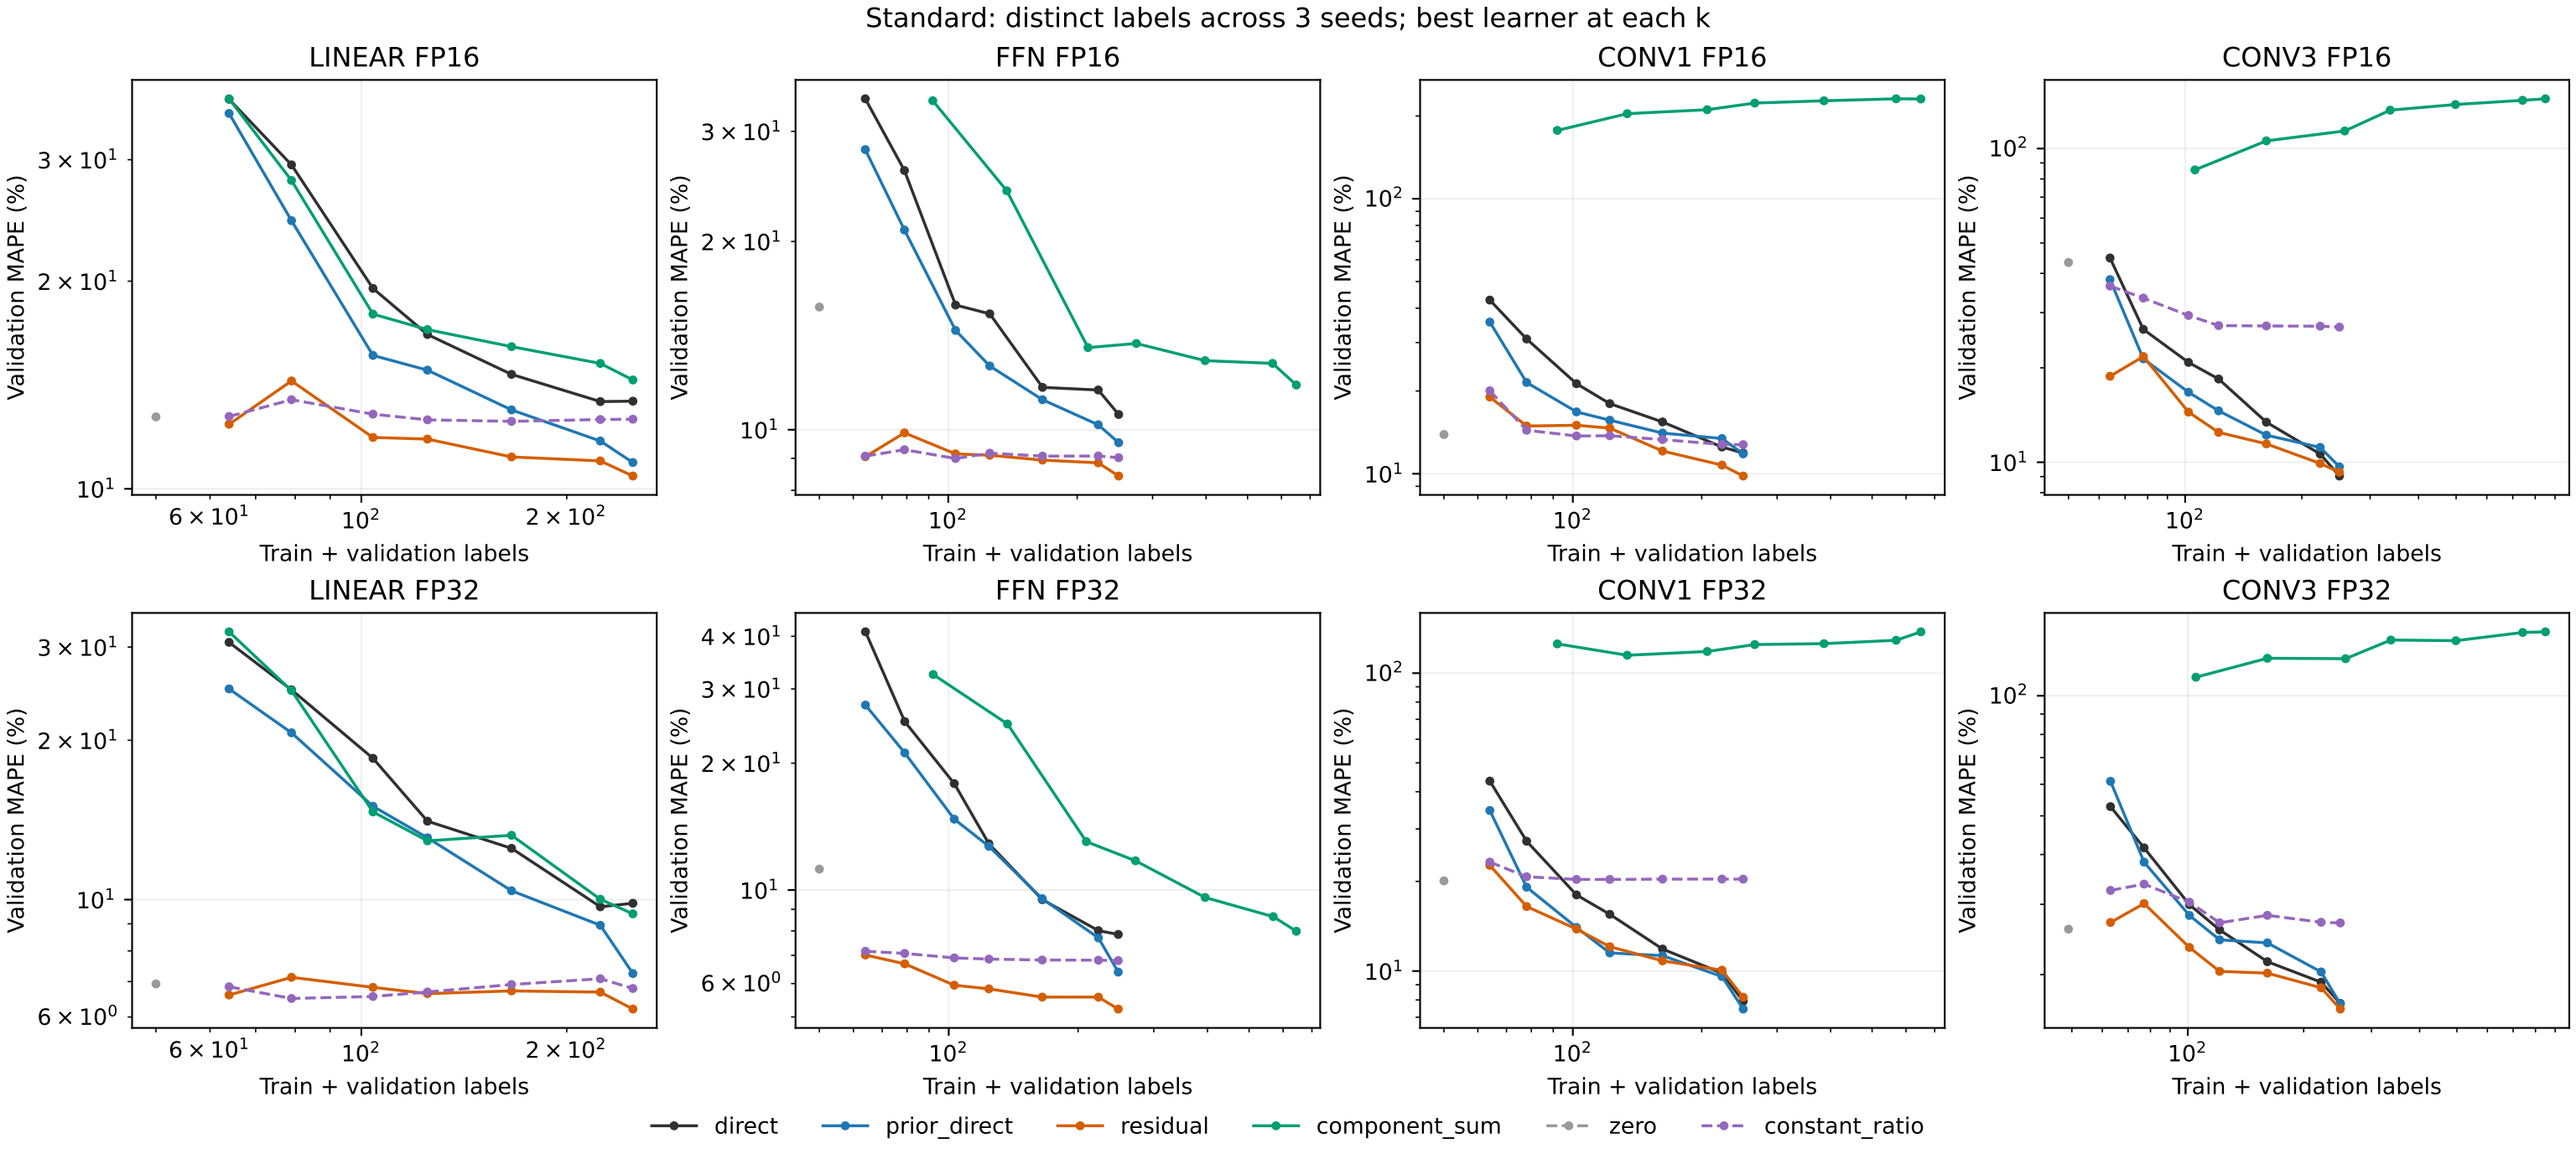
Standard: distinct labels across 3 seeds; best learner at each k
LINEAR FP16
FFN FP16
CONV1 FP16
CONV3 FP16
LINEAR FP32
FFN FP32
CONV1 FP32
CONV3 FP32
Validation MAPE (%)
Train + validation labels
direct
prior_direct
residual
component_sum
zero
constant_ratio